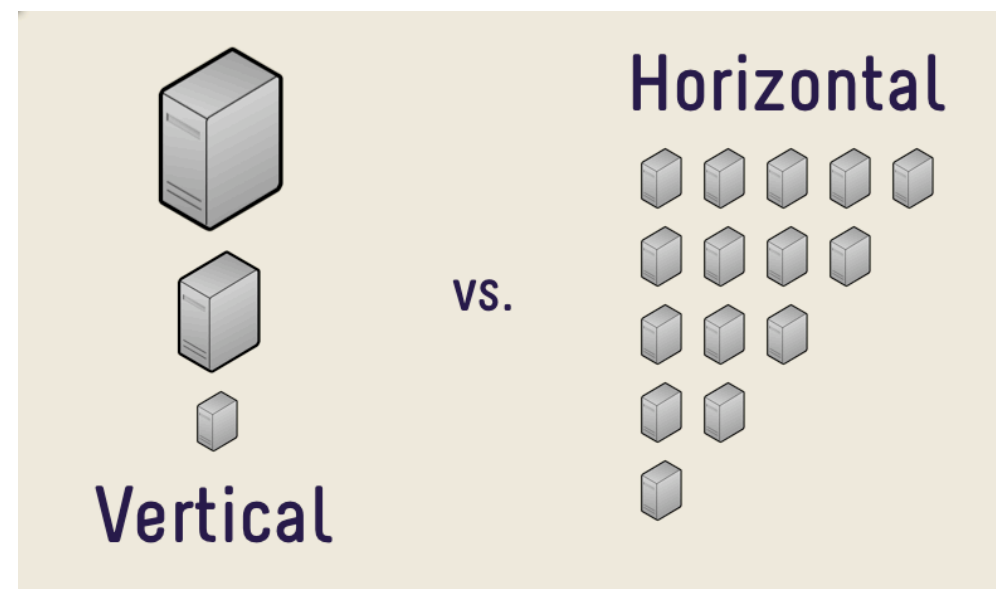


# Scalability

50.037 Blockchain Technology  
Paweł Szałachowski

# Scalability

- “*Scalability is the capability of a system, network, or process to handle a growing amount of work, or its potential to be enlarged to accommodate that growth.*”
- Horizontal and vertical scaling



# Blockchain Stack

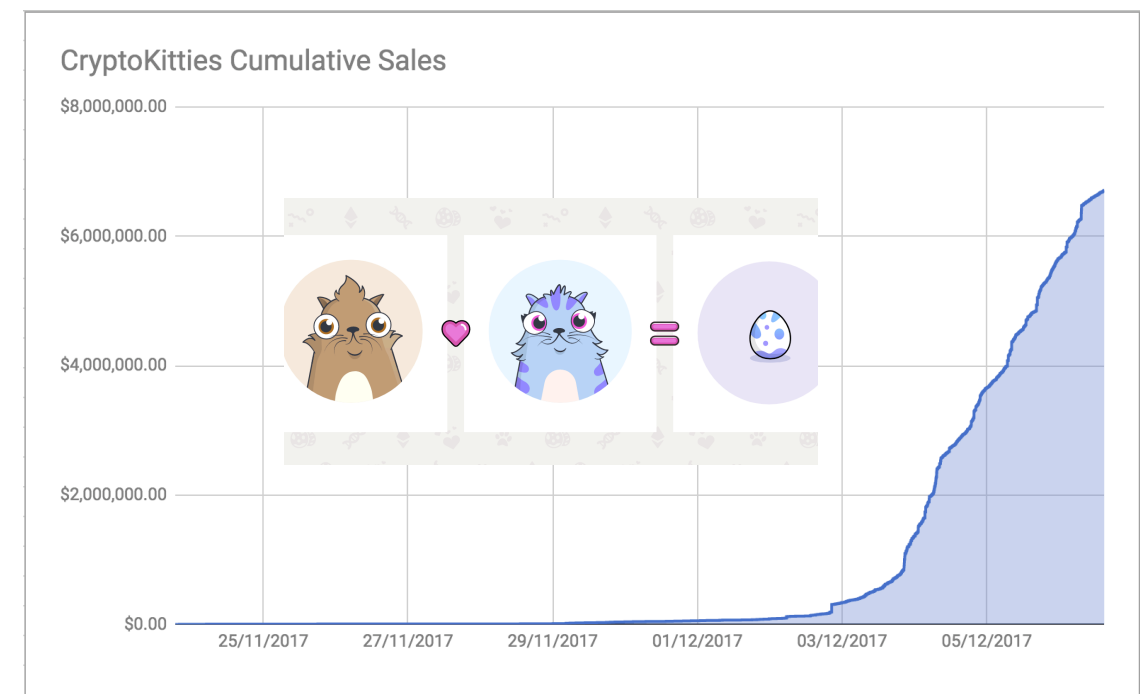
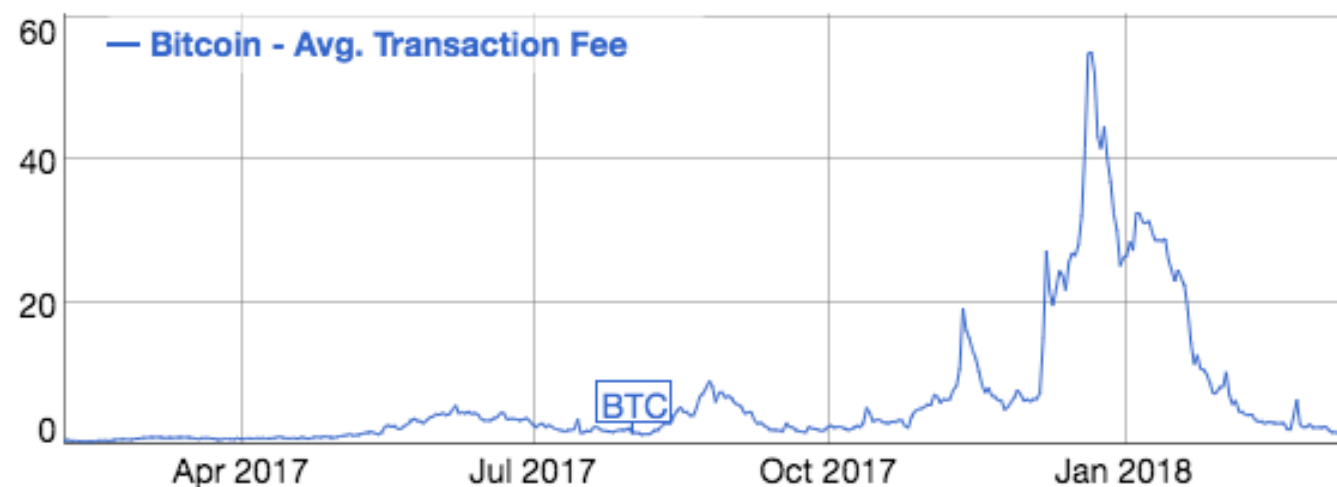
- Network: propagate transactions
  - Latency, bandwidth, # of nodes, ...
- **Consensus:** order Tx's
  - # of nodes, # of Tx's (throughput)
- RSM: Tx's validation, contract execution, ...
  - State size, execution complexity, ...
- Apps: use the current state to implement some logic

# Current State

- Throughput
  - Bitcoin: 7 tx/s
  - Ethereum: 10 tx/s
  - Visa: 50k tx/s
- Strategy
  - Faster tx processing
  - Faster consensus
  - Parallel execution

# Why needed?

- Adoption
  - Real-world apps require high throughput and low latency
- Inverse scale effect
- Fees



# Blockchain Performance

- Bandwidth:
  - How many Txs can be processed?
- Latency
  - What is the consensus delay?
- Mining power utilization
  - The ratio between the mining power of the current chain and the mining power of the entire blockchain (describes stale block rate too), describes *security*
- Fairness
  - A miner should benefit from rewards proportionally to its mining power

# Naive Improvements

- Blocks not every 10 minutes, but e.g., every 10 seconds
  - More forks => less mining power utilization => weaker security
- Larger blocks (very controversial topic BTW)
  - It takes longer to propagate
    - More forks => ....
- Bitcoin -> Bitcoin Cash -> Bitcoin ABC vs Bitcoin SV

# Security vs Performance

- Seems like Nakamoto consensus has some inherent tradeoffs
  - Security vs performance tradeoff
- Does it have to be like that?
  - We cannot significantly increase block or make them very frequent
- Design space
  - Why we need PoW?
  - Does it have to be combined with transactions propagation?




# Bitcoin-NG

- <https://www.usenix.org/node/194907> (paper, slides, talk)
- Insights
  - In Bitcoin, leader election and transaction serialization is combined
  - Why do not try to decouple it?
    - Elect leader via PoW and let her commit transactions
    - (Different order than in Bitcoin)

# Bitcoin-NG

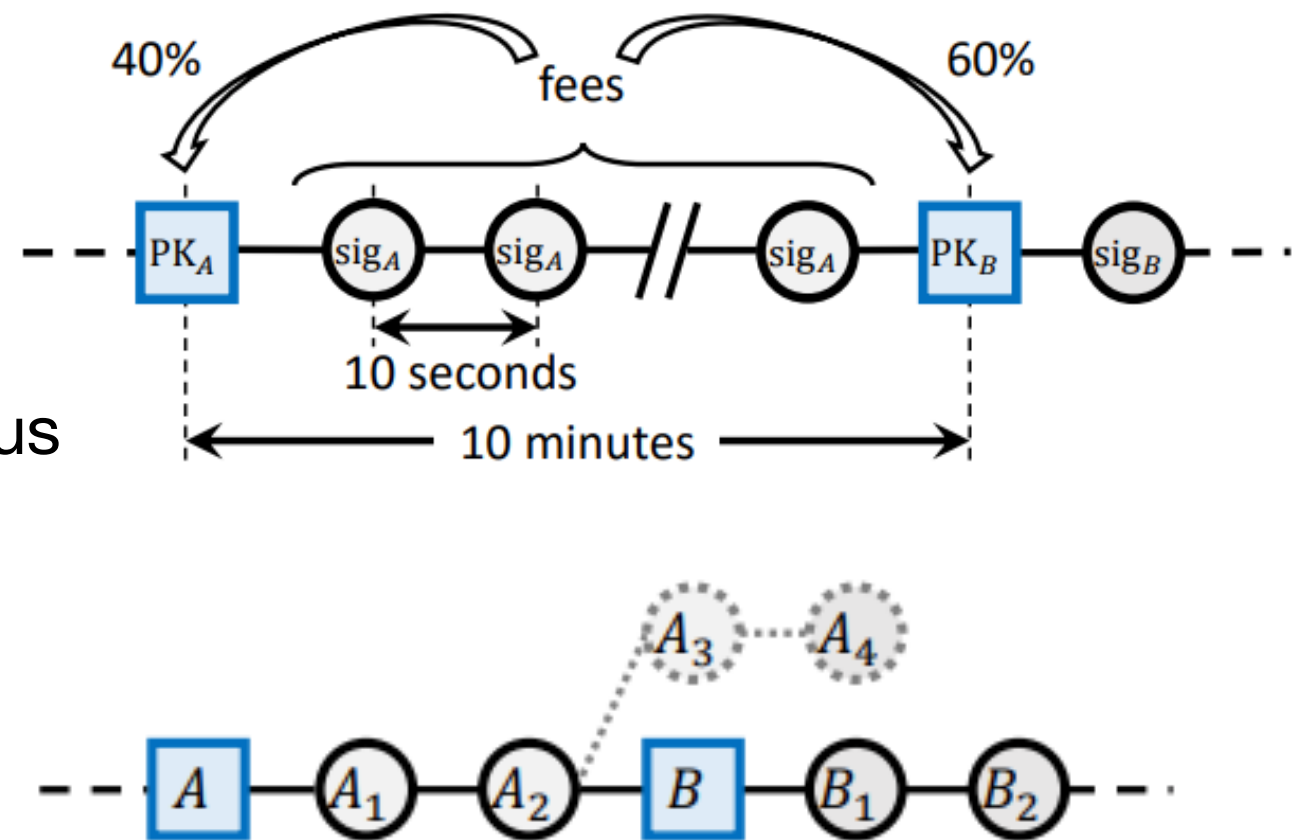
- Key blocks
  - Used for PoW-based leader election, i.e.,  $H(\text{header}) < T$
  - Point to the previous block (key or microblock)



- The strongest-chain rule
- Microblocks
  - Generated by leader, at a defined rate
  - Contains header (with PrevHash) and a set of transactions

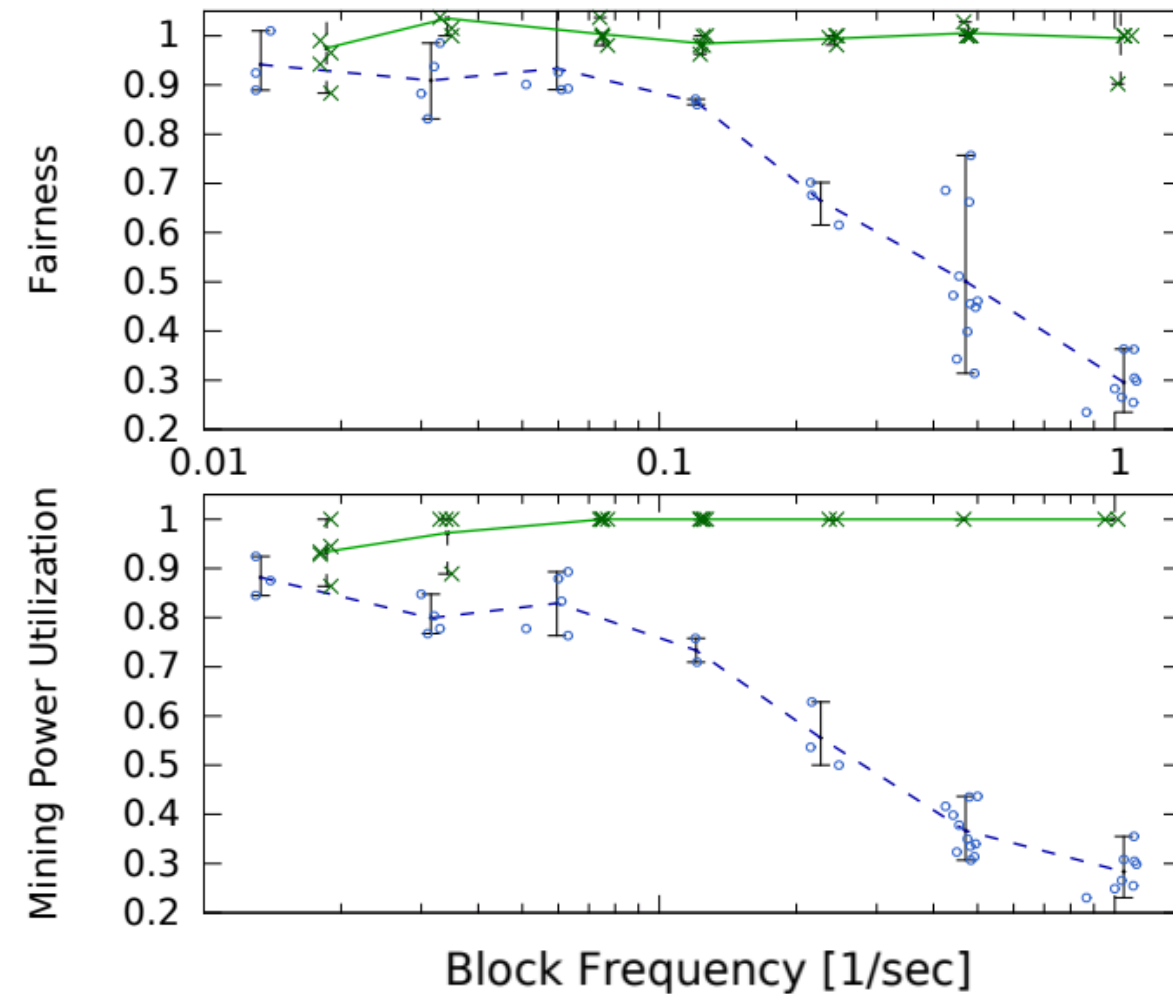
# Bitcoin-NG

- Incentives
  - Leaders get rewards and tx fees
  - The next leader gets 60% of the previous tx fees (why?)
- Confirmations
  - Short forks will be frequent
- Microblock forks may be malicious
  - Entry with a proof of fraud can invalidate the revenue of malicious leaders



# Bitcoin-NG

- Much better scalability and performance than Bitcoin
- Many systems build on this or similar ideas
- Everyone validates Tx's
  - Throughput limited by a single machine
- Can we do better?



# Sharding

# Sharding

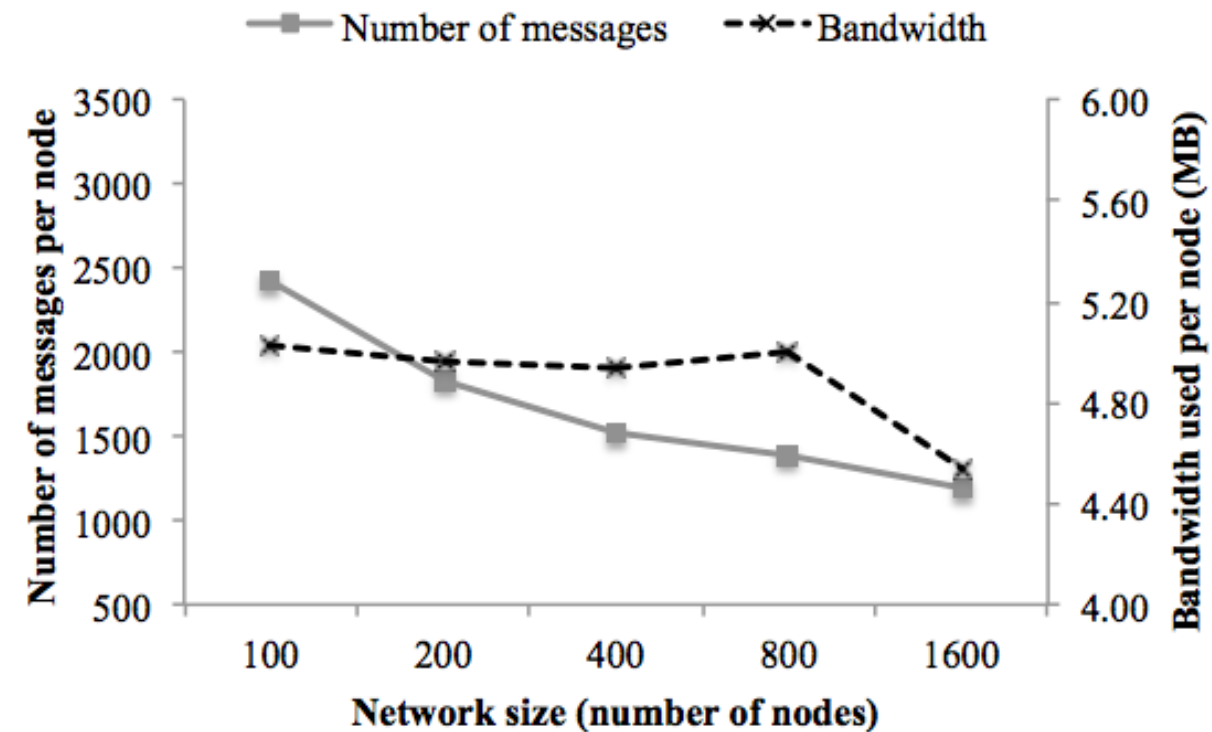
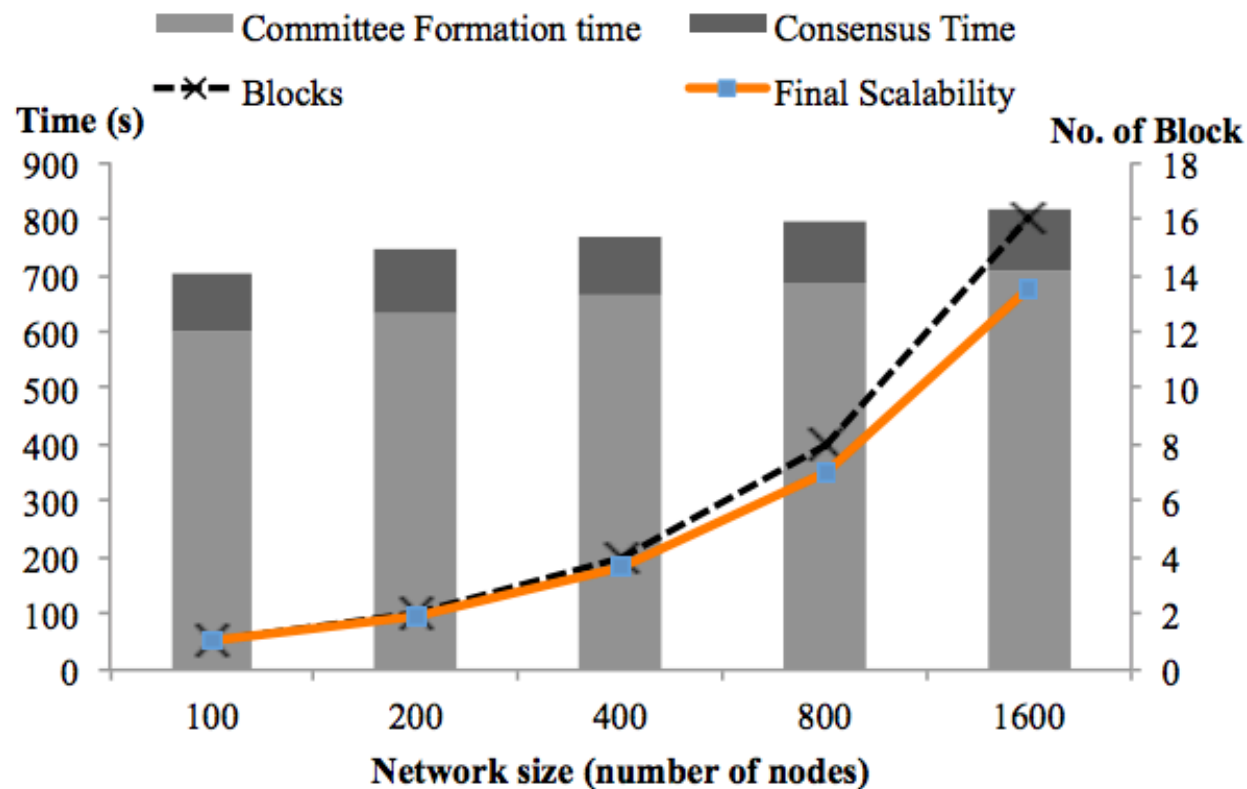
- The concept from database processing
- Divide transactions into groups and let different nodes process them
- Horizontal scaling
  - Throughput increases linearly as the network grows
- Ideas: establish identities via PoW, divide work, run BFT

# Sharding: Elastico

<https://dl.acm.org/citation.cfm?id=2978389>

1. Use PoW to establish identities
  - $ID = H(R, IP, PubKey, Nonce) < T$
  - R is security-critical, see below
2. Assign committees (use randomness of IDs)
  - Each committee has C members and a directory server (w/ members)
3. Propose a block within a committee
  - Run BFT agreement, valid blocks have  $2C/3 + 1$  signatures
4. Final committee to union all data blocks
  - Run BFT to produce a final block, that is then broadcast to everyone
  - R generated using  $R_x$  of final committee members

# Sharding: Elastico



- Multiple improvements (ongoing research)
- ZILLIQA, OmniLedger, Chainspace, Saber, ...



# Directed Acyclic Graph (DAG)

# DAG

- Nakamoto-like consensus is very easy
- The chain structure is simple
  - Can we introduce a more efficient data structure?
- Graphs
  - More powerful, as blocks can encode their worldviews
  - Why acyclic?

# SPECTRE

- <https://eprint.iacr.org/2016/1159.pdf> (payment oriented)
- Intuitions and insights:
  - DAG can be used to encode the longest-chain rule
  - Bitcoin is related to voting
    - Every block votes for its chain
    - Cloning in Bitcoin
  - Vote amplification (miners strengthen the majority decision)

# SPECTRE

- Miners create PoW-protected blocks with Txs
  - Blocks point to all known *tips* of the DAG
- Blocks can have conflicting Txs
- Nodes maintain local copies of DAG and accept/reject Tx
  - RobustReject, Pending, RobustAccepted

# SPECTRE

- Vote over blocks (pairwise)
  - How many think  $A < B$  vs how many think  $B < A$ 
    - This is just interpretation, does not have to be “*true*”
- Use results to accept/reject Tx's
  - Tx of block B if:
    - All inputs are accepted
    - For every conflicting Tx' in B',  $B < B'$

# SPECTRE





Fig. 1: An example of the voting procedure on a simple DAG. Block  $x$  and blocks 6-8 vote  $x \prec y$  as they only see  $x$  in their past, and not  $y$ . Similarly, block  $y$  and blocks 9-11 vote  $y \prec x$ . Block 12 votes according to a recursive call on the DAG that does not contain blocks 10,11,12. Any block from 1-5 votes  $x \prec y$ , because it sees more  $x \prec y$  voters in its future than  $y \prec x$  voters.

- Clone-proof and amplified decisions
- Quick confirmations

# Intel's Proof of Elapsed Time (PoET)

# Intel Software Guard Extensions (SGX)

- Set of new CPU instructions (Trusted Execution Environment — TEE)
- User code can allocate private regions of memory (**enclaves**), protected from other processes (even those running at higher privilege levels)
  - Running code and its memory is isolated from the rest of system
- Minimize trusted base
  - only CPU is trusted (even DRAM is untrusted, encryption needed)
- Attestation
  - Prove to remote system what code enclave is running





# PoET

- Observation: PoW is introduced to mitigate Sybil attacks
  - ... but with TEE that could be easier
- Idea: simulate PoW by sleep(...);
  - PoW-like guarantees
  - No energy waste
    - More energy vs more Intel SGX CPUs
- Intel & SGX are trusted, not fully open + see the recent attacks

# PoET

1. A newcomer node downloads the trusted code and sends a *join* message with the signed attestation
2. Nodes verify and accept/reject
3. In each round, every nodes gets a trusted random R and calls `sleep(R);`
4. The first awake node sends a signed msg that she is a leader
5. The statement is validated and the blocks can be produced

# Permissioned Blockchains

# Permissioned Blockchains

- We discuss open/permissionless blockchains so far
  - Great for some use cases, not so great for other
    - Good: Append-only, decentralized, available, robust, transparent...
    - Bad: Slow, expensive storage&computation, volatility, immature technology, publicly available data, difficult to manage/update...
      - Mainly caused by the permissionless setting
- Why not reuse some of those ideas and run classic BFT consensus?
- Efficiency could be one of the major benefits



# HYPERLEDGER


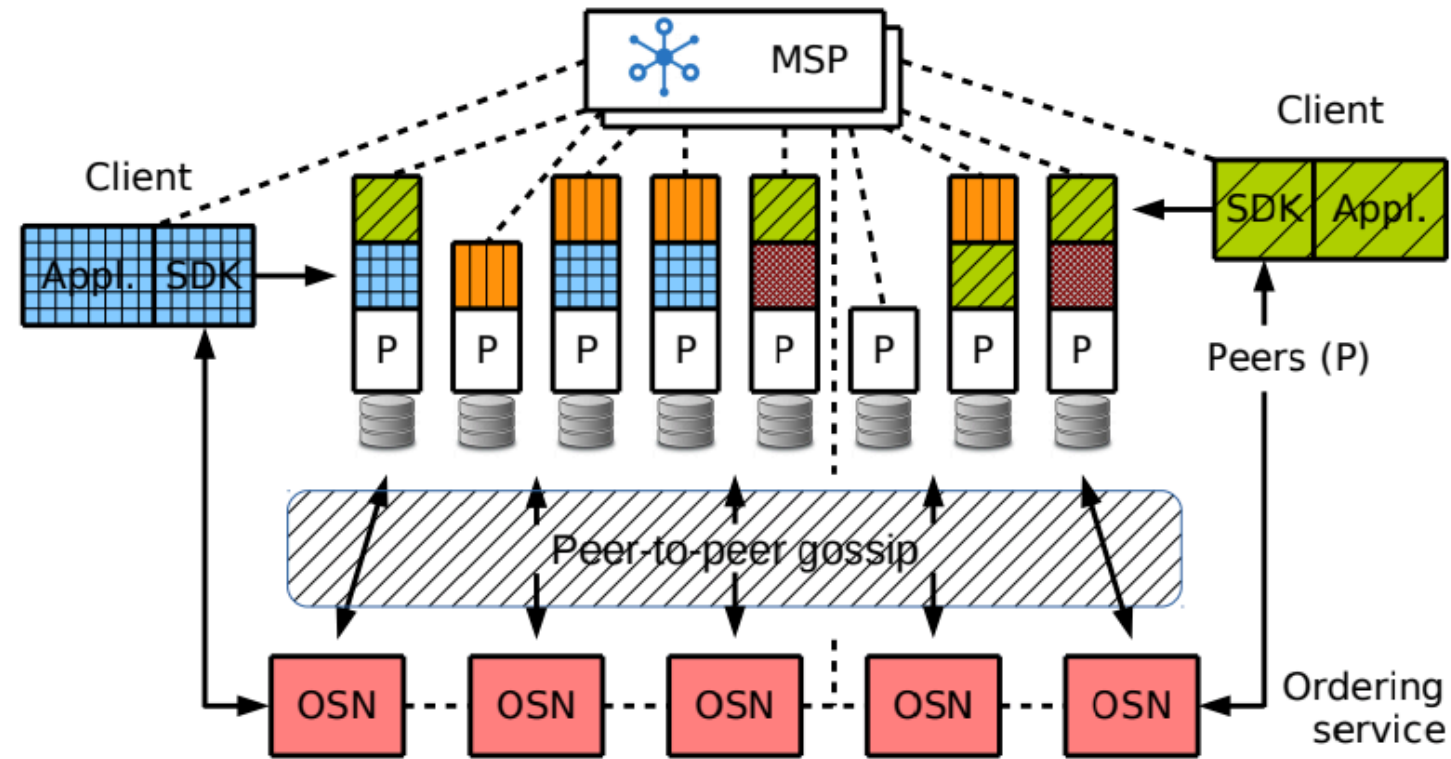
- Consortium
- Hyperledger Fabric (framework)
  - <https://arxiv.org/abs/1801.10228>
- Business-oriented
- Different consensus protocols supported
- No built-in cryptocurrency
- Powerful smart contracts



# Hyperledger Farbic

- Membership service provider (MSP) provides identities to participants
- Clients submit transaction proposals for execution
- Peers execute transaction proposals and validate transactions
  - Only endorsing peers execute transactions (specified by a policy)
  - All peers maintain the blockchain ledger
- Ordering Service Nodes (OSN) establish the total order of all transactions

# Hyperledger Fabric



# Reading

- <https://fc16.ifca.ai/bitcoin/papers/CDE+16.pdf>
- <https://eprint.iacr.org/2016/555.pdf>
- + inline references